Dynamic Programming and Reinforcement Learning

Lecturer: *Emilie Kaufmann* *emilie.kaufmann@inria.fr*

Report Deadline: November 13, midnight

Note: Your report should be *short* and based on the answers to questions Q1-Q4

Report and code should be sent by e-mail to emilie.kaufmann@inria.fr (one pdf file *yourname_TP1.pdf* + one archive *yourname_TP1.zip*), with [MVA 2016] in the title.

1 The One-Site Tree Cutting Problem

We would like to formalize the *tree cutting* problem and compute the strategy which maximizes the revenue. A tree keeps growing over time with a rate which may depend on the weather and it stops when it reaches a certain maximum height. At the same time the tree may get a disease, in which case it dies and looses all its value. When the company decides to cut a tree, it gains an amount of money which is proportional to the height of the tree. Whenever a tree is cut (or it is dead), a new tree has to be planted with a fixed cost. Knowing that the one unit of money looses value over time, find the optimal cutting strategy.

1.1 A Bit More Formal Definition of the Environment

- State space: the (discrete!) height of the tree (constrained to a maximum height)

- Initial state: the height of the tree is set to one

- Action space either cut or not the tree

- Dynamics:

    - If no cut: the tree grows up to a maximum height by a number of units which depend on the (random!) weather. It may also (randomly!) get a disease.

    - If cut: a new tree is planted with an initial height of one unit.

- Reward:

    - If no cut: a fixed amount of maintenance cost

    - If cut: the value of each unit of wood times the height of the tree minus the cost of planting a new tree.

- Discount factor: we assume a bank interest rate $r = 0.05$, and so discount factor is set of $\gamma = 1/(1+r)$.

1.2 Work to do

1. Formalize the problem more precisely (some decisions are of course arbitrary, such as the influence of the weather on the growth) and implement two functions:

    (a) tree_sim which receives as input a state and an action and it returns the next state and the reward.

    (b) tree_MDP which returns the dynamics and the reward function (in suitable structures).

    Q1: Explicit the MDP and the parameter chosen to model the random effects.

    Note: You may choose to use the representation proposed in the code *mainTP1.m* available online: the dynamics are represented by the "growth" matrix (that you can customize) and the sick probability. In that case, you can look into the function *tree_MDP.m* that is given. For *tree_simu.m* you may need to be able to simulate from finite distributions given by a probability vector: for this purpose you can use the given code *simu.m*.

2. Policy evaluation: define an arbitrary policy and evaluate it in the initial state using one RL method (Monte-Carlo or TD(0)) and one dynamic programming method (matrix inversion or Bellman operator).

    - Q2: If V_n denotes the value function computed by the RL method based on n trajectories, chart $\overline{V_n}(x_0) - V^\pi(x_0)$, where x_0 is the initial state and V^π is the value function computed with DP.

$$V_n(x_0) = \frac{1}{n} \sum_{k=1}^{n} \left[\sum_{t=1}^{T_{\max}} \gamma^{t-1} R_t^{(k)}, \right]$$

    where $(R_t^{(k)})$ is the sequence of rewards obtained when simulating the k-th trajectory (using *tree_simu*).

3. Optimal policy:

    - Q3: Compute the optimal policy with the two dynamic programming method seen in class, Policy Iteration and Value Iteration.
      Recall that both VI and PI can be implemented using the Q-value function associated to a value function V, defined by
$$Q(x, a) = r(x, a) + \gamma \sum_{y \in \mathcal{X}} p(y|x, a) V(y).$$

    - Q4: For both methods, plot $||V^* - V_k||_\infty$ as a function of iteration k to compare the speed of convergence and discuss the relative merits of the two approaches.
      For Policy Iteration, $V_k = V^{\pi_k}$, where π_k is the policy obtained after k iterations.

In the next session, we will implement Q-Learning for this problem (i.e. learning the optimal policy with a Reinforcement Learning approach)

2 Going further

1. Study how the obtained results change when changing some of the parameters of the problem (initial height, cost of planting a new tree, gain in selling a tree, and so on).

2. Consider the case where we have two sites where we can grow trees. At each point in time, the decision is whether to cut a tree and which one and the state should consider both sites. Implement the extension or discuss how it could be implemented.

3. Propose a model (and test Q-learning on it) to solve the problem sketched here `http://stackoverflow.com/questions/8337417/markov-decision-process-value-iteration-how-does-it-work`